

Pramathesh Shukla

CSC 576

Date: 11/2/2021

Hybrid Model for detecting Lung Diseases

1) Libraries

```
In [18]: import numpy as np
import pandas as pd
from os import listdir

import matplotlib.pyplot as plt
from sklearn.feature_selection import chi2

import cv2
import numpy as np
import cv2
from keras.preprocessing import image
from sklearn.preprocessing import LabelBinarizer
from sklearn.model_selection import train_test_split
import pickle

from tensorflow.keras.layers import Sequential
from tensorflow.keras.layers import BatchNormalization, Conv2D, MaxPooling2D, MaxPool2D, Flatten, Dense, Activation, Dropout, SpatialDropout2D, ZeroPadding2D, Convolution2D
from tensorflow.keras.optimizers import Adam
from tensorflow.keras import backend as K
from tensorflow.keras.optimizers import SGD

import tensorflow as tf

from tensorflow.keras.utils import to_categorical
```

2) Loading dataset & Visualizing

```
In [19]: sdir = r'../CSC_576/COVID-19_Radiography_Database'

In [21]: def getCount(directory_root):
    root_list = listdir(directory_root)
    dic = {}
    for label_folder in food_list:
        count = 1
        items = listdir((directory_root)/(label_folder))
        for label in items:
            count+=1
            dic[label_folder] = count
    return dic_

In [22]: dic = getCount(sdir)
labels = list(dic.keys())
values = list(dic.values())
print(dic)

{'COVID': 3617, 'Lung_Opacity': 6013, 'Normal': 10193, 'Viral_Pneumonia': 1346}

In [23]: df = pd.DataFrame(columns=['Labels', 'Count'])

df['Labels'] = values
df['Count'] = labels
display(df)

textprops = {'fontsize': 15} # Font size of text in pie chart
fig, ax = plt.subplots(1, 1)
ax.pie(values, labels=labels, colors=colors, shadow=True, explode=(0.1, 0.0, 1.0), textprops=textprops, radius=1.4,
startangle=-30, autopct='%1.1f%%')
ax.axis('equal')
ax.set_title('Lungs Dataset', fontsize=15)
plt.show()

Labels Count
0 COVID 3617
1 Lung_Opacity 6013
2 Normal 10193
3 Viral_Pneumonia 1346

Lungs Dataset
Normal
48.2%
Lung_Opacity
29.4%
Viral_Pneumonia
6.4%
COVID
17.1%
```

```
In [24]: def visualizeData(directory_root):
    food_list = listdir(directory_root)
    plt.figure(figsize=(35,20))
    dic = {}
    count = 1
    for label_folder in food_list:
        items = listdir((directory_root)/(label_folder))
        for label in items[5:]:
            plt.subplot(4,5,count)
            image_directory = (directory_root)/(label_folder)/(label)
            img = cv2.imread(image_directory)
            image = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
            plt.xticks([])
            plt.yticks([])
            plt.imshow(image, cmap=plt.cm.bone)
            plt.xlabel(label_folder, fontsize=25)
            count+=1
            plt.show()

    visualizeData(sdir)
```

```
In [26]: default_image_size = tuple((256, 256))

In [27]: def convert_image_to_array(image_dir):
    try:
        image = cv2.imread(image_dir)
        if image is not None:
            image = cv2.resize(image, default_image_size)
            return image_to_array(image)
        else:
            return np.array([])
    except Exception as e:
        print("Error : ", e)
        return None

In [28]: image_list, label_list = [], []
count = 0
try:
    print("[INFO] Loading images ...")
    root_dir = listdir(sdir)
    for label_folder in root_dir:
        covid_folder_list = listdir((sdir)/(label_folder))
        print("Processing: ", label_folder)
        for image in covid_folder_list[1346]: # Execute Negative & Positive Folder
            count+=1
            image_directory = (sdir)/(label_folder)/(image)
            image_list.append(convert_image_to_array(image_directory))
            label_list.append(label_folder)
    print("[INFO] Image loading completed")
except Exception as e:
    print("Error : ", e)

[INFO] Loading images ...
Processing: COVID
Processing: Lung_Opacity
Processing: Normal
Processing: Viral_Pneumonia
[INFO] Image loading completed

In [29]: label_binarizer = LabelBinarizer()
image_labels = label_binarizer.fit_transform(label_list)
image_labels_2 = to_categorical(image_labels)
pickle.dump(label_binarizer, open('label_transform.pkl', 'wb'))
n_classes = len(label_binarizer.classes_)

In [30]: print(label_binarizer.classes_)

['COVID' 'Lung_Opacity' 'Normal' 'Viral_Pneumonia']

In [31]: np_image_list = np.array(image_list, dtype=np.float16) / 255.0

In [32]: print("[INFO] Splitting data to train, test")
x_train, x_test, y_train, y_test = train_test_split(np_image_list, image_labels, image_labels_2,
                                                    test_size=0.2, random_state=42, shuffle=True)

[INFO] Splitting data to train, test

In [33]: print(x_train.shape, y_train.shape, '\n', x_test.shape, y_test.shape)

(4380, 256, 256, 3) (4380, 4)
(1877, 256, 256, 3) (1877, 4)

In [34]: aug = ImageDataGenerator(
    rotation_range=0,
    width_shift_range=0,
    height_shift_range=0,
    shear_range=0,
    zoom_range=0,
    horizontal_flip=False,
    fill_mode="nearest")

In [35]: # Fit parameters from data
aug.fit(x_train)

In [36]: # Configure batch size and retrieve one batch of images
for x_batch, y_batch in aug.flow(x_train, y_train, batch_size=32):
    # Show 8 images
    plt.figure(figsize=(25,10))
    for i in range(0, 8):
        plt.subplot(28, 1, i+1)
        plt.imshow(x_batch[i])
    # Show the plot
    plt.show()
    break
```

```
In [37]: # Fit parameters from data
aug.fit(x_train)

In [38]: # Configure batch size and retrieve one batch of images
for x_batch, y_batch in aug.flow(x_train, y_train, batch_size=32):
    # Show 8 images
    plt.figure(figsize=(25,10))
    for i in range(0, 8):
        plt.subplot(28, 1, i+1)
        plt.imshow(x_batch[i])
    # Show the plot
    plt.show()
    break
```

```
In [39]: EPOCHS = 10000
BS = 32
width=256
height=256
depth=3

Model Custom (CNN)

In [51]: model_custom = Sequential()
inputShape = (height, width, depth)
channels = 1
if K.image_data_format() == "channels_first":
    inputShape = (depth, height, width)
    channels = 1
model_custom.add(Conv2D(32, (3, 3), padding="same", activation="relu", input_shape=inputShape, strides=(1,1)))
model_custom.add(MaxPooling2D(pool_size=(2, 2)))
model_custom.add(Conv2D(64, (3, 3), padding="same", activation="relu", strides=(1,1)))
model_custom.add(Conv2D(64, (3, 3), padding="same", activation="relu", strides=(1,1)))
model_custom.add(MaxPooling2D(pool_size=(2, 2)))
model_custom.add(Conv2D(128, (3, 3), padding="same", activation="relu", strides=(1,1)))
model_custom.add(Conv2D(128, (3, 3), padding="same", activation="relu", strides=(1,1)))
model_custom.add(MaxPooling2D(pool_size=(2, 2)))
model_custom.add(Conv2D(512, (3, 3), padding="same", activation="relu", strides=(1,1)))
model_custom.add(Conv2D(512, (3, 3), padding="same", activation="relu", strides=(1,1)))
model_custom.add(Conv2D(512, (3, 3), padding="same", activation="relu", strides=(1,1)))
model_custom.add(MaxPooling2D(pool_size=(2, 2)))
model_custom.add(Flatten())
model_custom.add(Dense(1024, activation="relu"))
model_custom.add(Dropout(0.5))
model_custom.add(Dense(4, activation="softmax"))

In [52]: model_custom.summary()

Model: 'sequential_2'
Layer (type) Output Shape Param #
-----
conv2d_29 (Conv2D) (None, 256, 256, 32) 896
max_pooling2d_10 (MaxPooling) (None, 85, 85, 32) 0
conv2d_30 (Conv2D) (None, 85, 85, 64) 18496
conv2d_31 (Conv2D) (None, 85, 85, 64) 36928
max_pooling2d_11 (MaxPooling) (None, 42, 42, 64) 0
conv2d_32 (Conv2D) (None, 42, 42, 128) 73856
conv2d_33 (Conv2D) (None, 42, 42, 128) 147584
max_pooling2d_12 (MaxPooling) (None, 21, 21, 128) 0
conv2d_34 (Conv2D) (None, 21, 21, 512) 590336
conv2d_35 (Conv2D) (None, 21, 21, 512) 2359808
conv2d_36 (Conv2D) (None, 21, 21, 512) 2359808
max_pooling2d_13 (MaxPooling) (None, 10, 10, 512) 0
flatten_2 (Flatten) (None, 51200) 0
dense_7 (Dense) (None, 1024) 52429824
dropout_3 (Dropout) (None, 1024) 0
dense_8 (Dense) (None, 4) 4180
Total params: 80,021,608
Trainable params: 80,021,636
Non-trainable params: 0
```

```
In [53]: opt = SGD(lr=INIT_LR)
model_custom.compile(loss='mean_squared_error',
optimizer=opt, metrics=['accuracy'])

In [54]: history = model_custom.fit_generator(aug.flow(x_train, y_train, batch_size=BS),
validation_data=(x_test, y_test),
steps_per_epoch=len(x_train) // BS,
epochs=EPOCHS,
verbose=1)
```

```
Epoch 1/100
134/134 [=====] - 374s 3s/step - loss: 0.1872 - accuracy: 0.2620 - val_loss: 0.1869 - val_accuracy: 0.2488
Epoch 2/100
134/134 [=====] - 353s 3s/step - loss: 0.1865 - accuracy: 0.2624 - val_loss: 0.1862 - val_accuracy: 0.2835
Epoch 3/100
134/134 [=====] - 375s 3s/step - loss: 0.1858 - accuracy: 0.4992 - val_loss: 0.1854 - val_accuracy: 0.5102
Epoch 4/100
134/134 [=====] - 379s 3s/step - loss: 0.1846 - accuracy: 0.4345 - val_loss: 0.1836 - val_accuracy: 0.4267
Epoch 5/100
134/134 [=====] - 375s 3s/step - loss: 0.1818 - accuracy: 0.4919 - val_loss: 0.1786 - val_accuracy: 0.5135
Epoch 6/100
134/134 [=====] - 372s 3s/step - loss: 0.1723 - accuracy: 0.4550 - val_loss: 0.1597 - val_accuracy: 0.4828
Epoch 7/100
134/134 [=====] - 379s 3s/step - loss: 0.1516 - accuracy: 0.5349 - val_loss: 0.1366 - val_accuracy: 0.5997
Epoch 8/100
134/134 [=====] - 372s 3s/step - loss: 0.1428 - accuracy: 0.5788 - val_loss: 0.1319 - val_accuracy: 0.6138
Epoch 9/100
134/134 [=====] - 361s 3s/step - loss: 0.1346 - accuracy: 0.5998 - val_loss: 0.1409 - val_accuracy: 0.5596
Epoch 10/100
134/134 [=====] - 379s 3s/step - loss: 0.1295 - accuracy: 0.6132 - val_loss: 0.1211 - val_accuracy: 0.6416
Epoch 11/100
134/134 [=====] - 378s 3s/step - loss: 0.1220 - accuracy: 0.6420 - val_loss: 0.1175 - val_accuracy: 0.6676
Epoch 12/100
134/134 [=====] - 392s 3s/step - loss: 0.1280 - accuracy: 0.6533 - val_loss: 0.1140 - val_accuracy: 0.6657
Epoch 13/100
134/134 [=====] - 379s 3s/step - loss: 0.1184 - accuracy: 0.6589 - val_loss: 0.1131 - val_accuracy: 0.6815
Epoch 14/100
134/134 [=====] - 382s 3s/step - loss: 0.1114 - accuracy: 0.6774 - val_loss: 0.1099 - val_accuracy: 0.6862
Epoch 15/100
134/134 [=====] - 372s 3s/step - loss: 0.1080 - accuracy: 0.6869 - val_loss: 0.1073 - val_accuracy: 0.6899
Epoch 16/100
134/134 [=====] - 377s 3s/step - loss: 0.1059 - accuracy: 0.6916 - val_loss: 0.1079 - val_accuracy: 0.7224
Epoch 17/100
134/134 [=====] - 373s 3s/step - loss: 0.1021 - accuracy: 0.6978 - val_loss: 0.1067 - val_accuracy: 0.7289
Epoch 18/100
134/134 [=====] - 361s 3s/step - loss: 0.0995 - accuracy: 0.7068 - val_loss: 0.0955 - val_accuracy: 0.7252
Epoch 19/100
134/134 [=====] - 399s 3s/step - loss: 0.0945 - accuracy: 0.7251 - val_loss: 0.0900 - val_accuracy: 0.7465
Epoch 20/100
134/134 [=====] - 356s 3s/step - loss: 0.0891 - accuracy: 0.7312 - val_loss: 0.0863 - val_accuracy: 0.6843
Epoch 21/100
134/134 [=====] - 355s 3s/step - loss: 0.0916 - accuracy: 0.7423 - val_loss: 0.0863 - val_accuracy: 0.7437
Epoch 22/100
134/134 [=====] - 367s 3s/step - loss: 0.0888 - accuracy: 0.7424 - val_loss: 0.0851 - val_accuracy: 0.7502
Epoch 23/100
134/134 [=====] - 374s 3s/step - loss: 0.0859 - accuracy: 0.7492 - val_loss: 0.0822 - val_accuracy: 0.7623
Epoch 24/100
134/134 [=====] - 382s 3s/step - loss: 0.1114 - accuracy: 0.6774 - val_loss: 0.1099 - val_accuracy: 0.6862
Epoch 25/100
134/134 [=====] - 372s 3s/step - loss: 0.1080 - accuracy: 0.6869 - val_loss: 0.1073 - val_accuracy: 0.6899
Epoch 26/100
134/134 [=====] - 377s 3s/step - loss: 0.1059 - accuracy: 0.6916 - val_loss: 0.1079 - val_accuracy: 0.7224
Epoch 27/100
134/134 [=====] - 373s 3s/step - loss: 0.1021 - accuracy: 0.6978 - val_loss: 0.1067 - val_accuracy: 0.7289
Epoch 28/100
134/134 [=====] - 361s 3s/step - loss: 0.0995 - accuracy: 0.7068 - val_loss: 0.0955 - val_accuracy: 0.7252
Epoch 29/100
134/134 [=====] - 399s 3s/step - loss: 0.0945 - accuracy: 0.7251 - val_loss: 0.0900 - val_accuracy: 0.7465
Epoch 30/100
134/134 [=====] - 356s 3s/step - loss: 0.0891 - accuracy: 0.7312 - val_loss: 0.0863 - val_accuracy: 0.6843
Epoch 31/100
134/134 [=====] - 355s 3s/step - loss: 0.0916 - accuracy: 0.7423 - val_loss: 0.0863 - val_accuracy: 0.7437
Epoch 32/100
134/134 [=====] - 367s 3s/step - loss: 0.0888 - accuracy: 0.7424 - val_loss: 0.0851 - val_accuracy: 0.7502
Epoch 33/100
134/134 [=====] - 374s 3s/step - loss: 0.0859 - accuracy: 0.7492 - val_loss: 0.0822 - val_accuracy: 0.7623
Epoch 34/100
134/134 [=====] - 382s 3s/step - loss: 0.1114 - accuracy: 0.6774 - val_loss: 0.1099 - val_accuracy: 0.6862
Epoch 35/100
134/134 [=====] - 372s 3s/step - loss: 0.1080 - accuracy: 0.6869 - val_loss: 0.1073 - val_accuracy: 0.6899
Epoch 36/100
134/134 [=====] - 377s 3s/step - loss: 0.1059 - accuracy: 0.6916 - val_loss: 0.1079 - val_accuracy: 0.7224
Epoch 37/100
134/134 [=====] - 373s 3s/step - loss: 0.1021 - accuracy: 0.6978 - val_loss: 0.1067 - val_accuracy: 0.7289
Epoch 38/100
134/134 [=====] - 361s 3s/step - loss: 0.0995 - accuracy: 0.7068 - val_loss: 0.0955 - val_accuracy: 0.7252
Epoch 39/100
134/134 [=====] - 399s 3s/step - loss: 0.0945 - accuracy: 0.7251 - val_loss: 0.0900 - val_accuracy: 0.7465
Epoch 40/100
134/134 [=====] - 356s 3s/step - loss: 0.0891 - accuracy: 0.7312 - val_loss: 0.0863 - val_accuracy: 0.6843
Epoch 41/100
134/134 [=====] - 355s 3s/step - loss: 0.0916 - accuracy: 0.7423 - val_loss: 0.0863 - val_accuracy: 0.7437
Epoch 42/100
134/134 [=====] - 367s 3s/step - loss: 0.0888 - accuracy: 0.7424 - val_loss: 0.0851 - val_accuracy: 0.7502
Epoch 43/100
134/134 [=====] - 374s 3s/step - loss: 0.0859 - accuracy: 0.7492 - val_loss: 0.0822 - val_accuracy: 0.7623
Epoch 44/100
134/134 [=====] - 382s 3s/step - loss: 0.1114 - accuracy: 0.6774 - val_loss: 0.1099 - val_accuracy: 0.6862
Epoch 45/100
134/134 [=====] - 372s 3s/step - loss: 0.1080 - accuracy: 0.6869 - val_loss: 0.1073 - val_accuracy: 0.6899
Epoch 46/100
134/134 [=====] - 377s 3s/step - loss: 0.1059 - accuracy: 0.6916 - val_loss: 0.1079 - val_accuracy: 0.7224
Epoch 47/100
134/134 [=====] - 373s 3s/step - loss: 0.1021 - accuracy: 0.6978 - val_loss: 0.1067 - val_accuracy: 0.7289
Epoch 48/100
134/134 [=====] - 361s 3s/step - loss: 0.0995 - accuracy: 0.7068 - val_loss: 0.0955 - val_accuracy: 0.7252
Epoch 49/100
134/134 [=====] - 399s 3s/step - loss: 0.0945 - accuracy: 0.7251 - val_loss: 0.0900 - val_accuracy: 0.7465
Epoch 50/100
134/134 [=====] - 356s 3s/step - loss: 0.0891 - accuracy: 0.7312 - val_loss: 0.0863 - val_accuracy: 0.6843
Epoch 51/100
134/134 [=====] - 355s 3s/step - loss: 0.0916 - accuracy: 0.7423 - val_loss: 0.0863 - val_accuracy: 0.7437
Epoch 52/100
134/134 [=====] - 367s 3s/step - loss: 0.0888 - accuracy: 0.7424 - val_loss: 0.0851 - val_accuracy: 0.7502
Epoch 53/100
134/134 [=====] - 374s 3s/step - loss: 0.0859 - accuracy: 0.7492 - val_loss: 0.0822 - val_accuracy: 0.7623
Epoch 54/100
134/134 [=====] - 382s 3s/step - loss: 0.1114 - accuracy: 0.6774 - val_loss: 0.1099 - val_accuracy: 0.6862
Epoch 55/100
134/134 [=====] - 372s 3s/step - loss: 0.1080 - accuracy: 0.6869 - val_loss: 0.1073 - val_accuracy: 0.6899
Epoch 56/100
134/134 [=====] - 377s 3s/step - loss: 0.1059 - accuracy: 0.6916 - val_loss: 0.1079 - val_accuracy: 0.7224
Epoch 57/100
134/134 [=====] - 373s 3s/step - loss: 0.1021 - accuracy: 0.6978 - val_loss: 0.1067 - val_accuracy: 0.7289
Epoch 58/100
134/134 [=====] - 361s 3s/step - loss: 0.0995 - accuracy: 0.7068 - val_loss: 0.0955 - val_accuracy: 0.7252
Epoch 59/100
134/134 [=====] - 399s 3s/step - loss: 0.0945 - accuracy: 0.7251 - val_loss: 0.0900 - val_accuracy: 0.7465
Epoch 60/100
134/134 [=====] - 356s 3s/step - loss: 0.0891 - accuracy: 0.7312 - val_loss: 0.0863 - val_accuracy: 0.6843
Epoch 61/100
134/134 [=====] - 355s 3s/step - loss: 0.0916 - accuracy: 0.7423 - val_loss: 0.0863 - val_accuracy: 0.7437
Epoch 62/100
134/134 [=====] - 367s 3s/step - loss: 0.0888 - accuracy: 0.7424 - val_loss: 0.0851 - val_accuracy: 0.7502
Epoch 63/100
134/134 [=====] - 374s 3s/step - loss: 0.0859 - accuracy: 0.7492 - val_loss: 0.0822 - val_accuracy: 0.7623
Epoch 64/100
134/134 [=====] - 382s 3s/step - loss: 0.1114 - accuracy: 0.6774 - val_loss: 0.1099 - val_accuracy: 0.6862
Epoch 65/100
134/134 [=====] - 372s 3s/step - loss: 0.1080 - accuracy: 0.6869 - val_loss: 0.1073 - val_accuracy: 0.6899
Epoch 66/100
134/134 [=====] - 377s 3s/step - loss: 0.1059 - accuracy: 0.6916 - val_loss: 0.1079 - val_accuracy: 0.7224
Epoch 67/100
134/134 [=====] - 373s 3s/step - loss: 0.1021 - accuracy: 0.6978 - val_loss: 0.1067 - val_accuracy: 0.7289
Epoch 68/100
134/134 [=====] - 361s 3s/step - loss: 0.0995 - accuracy: 0.7068 - val_loss: 0.0955 - val_accuracy: 0.7252
Epoch 69/100
134/134 [=====] - 399s 3s/step - loss: 0.0945 - accuracy: 0.7251 - val_loss: 0.0900 - val_accuracy: 0.7465
Epoch 70/100
134/134 [=====] - 356s 3s/step - loss: 0.0891 - accuracy: 0.7312 - val_loss: 0.0863 - val_accuracy: 0.6843
Epoch 71/100
134/134 [=====] - 355s 3s/step - loss: 0.0916 - accuracy: 0.7423 - val_loss: 0.0863 - val_accuracy: 0.7437
Epoch 72/100
134/134 [=====] - 367s 3s/step - loss: 0.0888 - accuracy: 0.7424 - val_loss: 0.0851 - val_accuracy: 0.7502
Epoch 73/100
134/134 [=====] - 374s 3s/step - loss: 0.0859 - accuracy: 0.7492 - val_loss: 0.0822 - val_accuracy: 0.7623
Epoch 74/100
134/134 [=====] - 382s 3s/step - loss: 0.1114 - accuracy: 0.6774 - val_loss: 0.1099 - val_accuracy: 0.6862
Epoch 75/100
134/134 [=====] - 372s 3s/step - loss: 0.1080 - accuracy: 0.6869 - val_loss: 0.1073 - val_accuracy: 0.6899
Epoch 76/100
134/134 [=====] - 377s 3s/step - loss: 0.1059 - accuracy: 0.6916 - val_loss: 0.1079 - val_accuracy: 0.7224
Epoch 77/100
134/134 [=====] - 373s 3s/step - loss: 0.1021 - accuracy: 0.6978 - val_loss: 0.1067 - val_accuracy: 0.7289
Epoch 78/100
134/134 [=====] - 361s 3s/step - loss: 0.0995 - accuracy: 0.7068 - val_loss: 0.0955 - val_accuracy: 0.7252
Epoch 79/100
134/134 [=====] - 399s 3s/step - loss: 0.0945 - accuracy: 0.7251 - val_loss: 0.0900 - val_accuracy: 0.7465
Epoch 80/100
134/134 [=====] - 356s 3s/step - loss: 0.0891 - accuracy: 0.7312 - val_loss: 0.0863 - val_accuracy: 0.6843
Epoch 81/100
134/134 [=====] - 355s 3s/step - loss: 0.0916 - accuracy: 0.7423 - val_loss: 0.0863 - val_accuracy: 0.7437
Epoch 82/100
134/134 [=====] - 367s 3s/step - loss: 0.0888 - accuracy: 0.7424 - val_loss: 0.0851 - val_accuracy: 0.7502
Epoch 83/100
134/134 [=====] - 374s 3s/step - loss: 0.0859 - accuracy: 0.7492 - val_loss: 0.0822 - val_accuracy: 0.7623
Epoch 84/100
134/134 [=====] - 382s 3s/step - loss: 0.1114 - accuracy: 0.6774 - val_loss: 0.1099 - val_accuracy: 0.6862
Epoch 85/100
134/134 [=====] - 372s 3s/step - loss: 0.1080 - accuracy: 0.6869 - val_loss: 0.1073 - val_accuracy: 0.6899
Epoch 86/100
134/134 [=====] - 377s 3s/step - loss: 0.1059 - accuracy: 0.6916 - val_loss: 0.1079 - val_accuracy: 0.7224
Epoch 87/100
134/134 [=====] - 373s 3s/step - loss: 0.1021 - accuracy: 0.6978 - val_loss: 0.1067 - val_accuracy: 0.7289
Epoch 88/100
134/134 [=====] - 361s 3s/step - loss: 0.0995 - accuracy: 0.7068 - val_loss: 0.0955 - val_accuracy: 0.7252
Epoch 89/100
134/134 [=====] - 399s 3s/step - loss: 0.0945 - accuracy: 0.7251 - val_loss: 0.0900 - val_accuracy: 0.7465
Epoch 90/100
134/134 [=====] - 356s 3s/step - loss: 0.0891 - accuracy: 0.7312 - val_loss: 0.0863 - val_accuracy: 0.6843
Epoch 91/100
134/134 [=====] - 355s 3s/step - loss: 0.0916 - accuracy: 0.7423 - val_loss: 0.0863 - val_accuracy: 0.7437
Epoch 92/100
134/134 [=====] - 367s 3s/step - loss: 0.0888 - accuracy: 0.7424 - val_loss: 0.0851 - val_accuracy: 0.7502
Epoch 93/100
134/134 [=====] - 374s 3s/step - loss: 0.0859 - accuracy: 0.7492 - val_loss: 0.0822 - val_accuracy: 0.7623
Epoch 94/100
134/134 [=====] - 382s 3s/step - loss: 0.1114 - accuracy: 0.6774 - val_loss: 0.1099 - val_accuracy: 0.6862
Epoch 95/100
134/134 [=====] - 372s 3s/step - loss: 0.1080 - accuracy: 0.6869 - val_loss: 0.1073 - val_accuracy: 0.6899
Epoch 96/100
134/134 [=====] - 377s 3s/step - loss: 0.1059 - accuracy: 0.6916 - val_loss: 0.1079 - val_accuracy: 0.7224
Epoch 97/100
134/134 [=====] - 373s 3s/step - loss: 0.1021 - accuracy: 0.6978 - val_loss: 0.1067 - val_accuracy: 0.7289
Epoch 98/100
134/134 [=====] - 361s 3s/step - loss: 0.0995 - accuracy: 0.7068 - val_loss: 0.0955 - val_accuracy: 0.7252
Epoch 99/100
134/134 [=====] - 399s 3s/step - loss: 0.0945 - accuracy: 0.7251 - val_loss: 0.0900 - val_accuracy: 0.7465
Epoch 100/100
134/134 [=====] - 356s 3s/step - loss: 0.0891 - accuracy: 0.7312 - val_loss: 0.0863 - val_accuracy: 0.6843
```

```
In [56]: print("[INFO] Calculating Custom Model Accuracy")
scores = model_custom.evaluate(x_test, y_test)
print("Test Accuracy: (scores[1]*100)")

[INFO] Calculating Custom Model Accuracy
34/34 [=====] - 22s 637ms/step - loss: 0.0608 - accuracy: 0.8422
Test Accuracy: 84.2154545467837

In [57]: acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs_range = range(1, len(history.epoch) + 1)

plt.figure(figsize=(10,5))
plt.plot(epochs_range, acc, label='Training')
plt.plot(epochs_range, val_acc, label='Validation')
plt.legend(loc='best')
plt.xlabel('Epochs')
plt.ylabel('Model Accuracy')
plt.grid(True, which='major', color='#666666', linestyle='-')
plt.tight_layout()
plt.show()

plt.figure(figsize=(10,5))
plt.plot(epochs_range, loss, label='Training')
plt.plot(epochs_range, val_loss, label='Validation')
plt.legend(loc='best')
plt.xlabel('Epochs')
plt.ylabel('Model Loss')
plt.grid(True, which='major', color='#666666', linestyle='-')
plt.tight_layout()
plt.show()

Model Accuracy
Accuracy
0.9
0.8
0.7
0.6
0.5
0.4
0.3
0
0 20 40 60 80 100
Epochs
Training
Validation

Model Loss
Loss
0.175
0.150
0.125
0.100
0.075
0.050
0.025
0.000
0
20 40 60 80 100
Epochs
Training
Validation
```

```
In [58]: print("[INFO] Saving Custom model H5...")
model_custom.save(r'./Models/Custom.h5')
print("Model saved")

[INFO] Saving Custom model H5...
Model saved
```